

P0Pxp[;4YA+~@
l|S0Ru[;086E,~i0:Ahc/>*c>q*8;âB5|,IqUwû[Y9=î-u;~J"zçZž0*81/Uz;\fTW\R0Esi:.Ç"8uh[ZABOjectA
äv?_äZvD±g(YÜ"çwëo(ô|iY,qj
öŽe'Äg(f)Næ[U9;08uf12]0C;"çm7*-QzTce02y6'an:JOiei»YÁPUG*x+Us¥(%Å),(fUüPt3o/[PIQ:-W·ÜStæ>ôÖxÔÿÊA
uc*EHGJw-m)B78-W7ç@æfXkw'i0M-6|8Z:*Xè:â0)*I-TNâAmX/oë'á-EIOæ»pXššxp,Iç=-ôieÜQsv0šape(ifim*FAÜ|f0At;-uy
~Yi[;5XeiZ

AAYUSH GUPTA

77

General instructions:

1. Processes and related abstractions

- In Linux, there are two separate stacks for each process - a *user mode* stack and a *kernel mode* stack. Why are there two stacks?
- In Linux, the kernel mode stack is maintained within the same `struct` that holds the PCB. What benefit(s) does this offer?
- Linux reserves only 8KB for the kernel mode stack. What do you think happens when this stack overflows?

Kernel mode stack stores information regarding the state of execution of a thread when it goes from user-mode to kernel mode. It also contains scheduling & accounting information used by the kernel. Whenever an interrupt occurs the thread switches from user-mode to kernel mode & the kernel mode stack is used to set up user-mode stack for signal handler & perform signal handling functions like `setjmp()` etc.

The question WHY was not what these stacks do.

User mode stack contains all the user defined variables whereas when the thread runs in kernel-mode it uses the thread's kernel mode stack & helps in faster kernel-mode processing

b)

- Easier management as it is easier to attribute a kernel mode stack of a process if its in the same place as its PCB

OK

- The kernel can easily locate the kernel mode stack of a process in $O(1)$ time & also prevents duplication of features (reduces redundancy)

Why would there be a duplication otherwise?

More benefits - See solutions

Processes, etc.

- (c) Each thread is allocated a kernel mode stack so increasing the size of the stack by even a small amount will drastically increase the memory footprint of the system.
- It may crash if the kernel stack overflows. This might happen if recursive calls are stored in kernel mode stack.

5/5

1 $\frac{S_A}{5}$
2 $\frac{5}{2} \frac{5}{2}$

2. CPU scheduling

(10+5=15 pts)


- (a) Consider a computer in which the hardware timer starts malfunctioning in the following manner. Instead of interrupting the CPU once every 10 time units (which is what the OS expects), it starts interrupting it every 5 time units (without the OS realizing this.) What problems is it likely to create for the CPU scheduler, if it is: (i) a proportionate-fair scheduler or (ii) a reservation-based scheduler? What would happen if the malfunction results in interrupts every 20 time units instead of every 10 time units?
- (b) Consider the SFQ scheduler. How will the behavior of the scheduler change if we started scheduling processes in an increasing order of their finish tags (instead of their start tags)?

(a) (i) Proportionate-fair scheduler \rightarrow The proportionate fair scheduler allocates CPU to ^{runnable} processes in the ratio of their weights. Even though the CPU gets interrupted after 5 time units instead of 10, the proportionate-fair scheduler will be able to maintain fairness of CPU distribution amongst the processes. The only impact will be that a scheduling decision will be made every 5 time units instead of 10 (if the running process doesn't block) though this will remain hidden from the scheduler.

eg Thread A
($W=1$)

Actual Distribution

Thread B
($W=2$)



The proportion is maintained though the absolute amount per scheduling decision is decreased i.e. instead of 10 units A gets scheduled out after 5 units


If interrupts occur after 20 time units, the proportion will still be maintained (fairness), the actual allocated time will increase.

CPU scheduling ...


(ii) Reservation Based Scheduler \rightarrow The RB scheduler gives absolute guarantees & is non work conserving.
eg: $A(x, y) \rightarrow$ The process A is guaranteed X amount of CPU time every Y time units

In this case too the process will get X amount of CPU time every Y units though the scheduling will occur differently.

eg: $A(10, 20) \rightarrow$ The RB scheduler guaranteed 10 ^{time} units worth CPU every 20 time units in CPU available to A



Similarly even if the interrupt occurs after 20 time units, the guarantees would be met though in the above stated manner.
Continuing the same example



CPU scheduling ...

(b) SFO scheduler will degenerate to WFQ scheduler.

To compute finish tag the length of the quantum needs to be known a priori and scheduling according to finish tags will imply that this quantum be known beforehand. If for calculation max. quantum length is taken & the thread uses less than max. then it will not receive its fair share

Good, $\left(\frac{5}{8} \right)$

3. Memory management


(5+10+15=30 pts)

- (a) Discuss why Linux reserves a fixed portion of every address space for kernel addresses. Are there any downsides to this choice?
- (b) Consider an operating system that implements paging-based virtual memory management. Suppose you wrote a program that malloc-ed an int and then looped infinitely. Suppose we started a new copy of this program every minute. Disregarding memory allocated for anything other than for these processes, how would the memory consumption of the computer evolve with time?
- (c) Consider a typical multi-programming system with a buffer cache that the OS maintains for improving the efficiency of disk I/O. We want to implement a predictable virtual memory manager (with a well-defined objective such as: "maximize overall system throughput" or "share page fault rates proportionally among the processes"). One way of doing this is to maintain information that lets us estimate working sets of processes (such as the gLRU QPosition versus # accesses histogram) and use it to partition RAM among the processes. In this question, we would like to see your ideas on including the buffer cache into the predictable RAM partitioning problem.
- Here are some high-level proposals. For each proposal, explain what you understand it to be and any thoughts you may have on how to implement it. (Extra credit: You are welcome to make your own proposals, if you have any.) Finally, compare these proposals with each other.
- *Proposal 0:* Keep a fixed number of frames reserved for the buffer cache. Use gLRU within the buffer cache.
 - *Proposal 1:* Treat the buffer cache as RAM allocated to a separate (non-existent) process. Use gLRU within the buffer cache.
 - *Proposal 2:* View blocks in the buffer cache as part of the address space of processes they belong to.

(a) Reserving a fixed portion of every address space for kernel addresses prevents TLB flush. Whenever the process changes from user mode to kernel mode & generates a particular virtual address of kernel no TLB flush occurs.

5/5 The downside is that this kernel address space is fixed & if the kernel addresses have to increase then either the boundary has to be made dynamic i.e. user space has to be eaten up or the kernel might crash.

- (b) The memory consumption will increase with each malloc in a staircase manner




The first malloc will get a chunk of memory & will allocate "int" as long as it's available. When it's finished it will get another chunk & so on. Since there are no "free" calls and the programs keep looping thus the memory never gets released.

LOW \Rightarrow no actual RAM allocated!

~~What will happen when the virtual address space gets exhausted?~~
Till when does this continue?

Memory management ...

Proposal 0 →
(C) Similar to segmented queue
use GURU here



Proposal 1 → The memory manager will find it as another process & assign memory according to its objective. The process here can then implement its own memory manager, GURU etc. (like Memesin)

Proposal 2 → This amounts to extending the address space of a process. The portions of buffer cache become part of the process & then will be treated as other process pages.

how? What if buffer cache is almost full?
Proposal 0 will improve disk I/O but may impact other processes as total RAM available to them will decrease. Proposal 1 may not improve I/O if all the pages for buffer cache get overwritten by other process pages. Proposal 2 gives processes the

advantage of managing the buffer cache allocated to them by themselves

4. I/O subsystem

(10+10+10=30 pts)

- (a) Consider a computer running several backlogged TCP and UDP connections. Suppose we wanted to implement a proportionate-share (PS) scheduling algorithm to isolate these connections from each other. Where in the OS would you implement this? Draw parallels with a CPU scheduling system that is proportionally sharing cycles among several CPU-intensive and I/O-intensive processes. (Hint: Are certain connections like CPU-intensive processes while others are like I/O-intensive processes? In what ways?) Can you also think of any significant differences between PS scheduling for CPU and network bandwidth?
- (b) Consider a computer with multiple Network Interface Cards (NICs.) Suppose we wanted to implement proportionate sharing (PS) of network bandwidth in such a system. Let us compare this with PS of CPU cycles in a multi-processor computer. What do you think are the counterparts of processes and processors in this setting? Do problems similar to "Weight Infeasibility" and "Short Jobs Arrival" that we studied for multi-processor CPU scheduling arise here? Explain your answer.
- (c) Can you think of a resource for which being non-work-conserving can actually improve throughput? Make sure you clearly state what aspect of the operation of this resource you consider to be non-work-conserving.

(a) UDP connections can be compared to CPU intensive processes whereas TCP connections are like I/O intensive processes because there may be fragmentation in TCP & a packet may have to wait for all fragments before it can be further processed

CPU PS scheduler does not give extra advantage to I/O intensive processes. Whenever a blocked process is ready to run, the scheduler gives it the start tag same as virtual time (min. of running threads' start tags)

The CPU scheduler¹⁰ can account resources to the process/thread which consumed them i.e. process is the right abstraction for resource

7/10

I/O subsystem ...

principal whereas in n/w scheduler the processing for a particular connection may involve multiple threads.

- (b) NIC cards represent processors
Each socket (fd) is a process.

Yes the same problems are there in PS of n/w bandwidth.

If the weights assigned to each socket (application) are infeasible i.e. more bandwidth is assigned than NIC capacity then "weight infeasibility" will occur.

The "short jobs arrival" is also there.

An Application (socket) which receives packets intermittently may get the same bandwidth as a higher priority (more weight) process.

7/10

I/O subsystem ...

- (C) Disk is the resource for which being non-conserving can improve throughput. Data is spread on a disk & it may not exhibit spatial locality ^{for a process}. A proportionate scheduler may try to access data which is far away from the head. This will be an expensive operation whereas a non-work conserving scheduler instead of moving the head to a farther position will keep the head idle to improve throughput.



5. Feedback

- (a) Write two things you like about CSE 511.
(b) Write two things you dislike about CSE 511.

(5+5=10 pts)

(a) Likes:

- (1) The best thing I like is not accepting things written in papers as it is. We try to challenge the notions & try to bring out contradictions & flaws in the given arguments & accept things only when totally convinced.
- (2) The open discussion oriented atmosphere and not imposing a particular point of view.

(b) Dislikes →

- (1) Sometimes there is a mismatch in what was covered in previous class & what happens in the current class. Probably the flow of things can be improved.

- (2) Certain things have remained undiscussed like scheduler activations.

It is the first time because of my first 511 for the first time.

Thanks. I wish you had given a concrete example. Did you mean "jumping across topics" or "contradiction between things taught in one class and next"?

If on the other hand... it is then that is a problem and you should tell me about when that happened.

10/10